

# CMPT 413

## Computational Linguistics

Anoop Sarkar

<http://www.cs.sfu.ca/~anoop>

2/20/08

1


## Sequence Learning

- British Left Waffles on Falkland Islands
  - (N, N, V, P, N, N)
  - (N, V, N, P, N, N)
- Segmentation 中国十四个边境开放城市经济建设成就显著
  - (b, i, b, i, b, b, i, b, i, b, i, b, i, b, i, b, i)
  - 中国 十 四 个 边 境 开 放 城 市 经 济 建 设 成 就 显 著
  - China 's 14 open border cities marked economic achievements

2/20/08

2

# Sequence Learning



**3 states:** N, V, P


**Observation sequence:**  $(o_1, \dots, o_6)$

**State sequence (6+1):**  $(Start, N, N, V, P, N, N)$

2/20/08

3

# Finite State Machines




Mealy Machine

2/20/08

4

# Finite State Machines



Moore Machine

2/20/08

5


## Probabilistic FSMs

- Each transition is associated with a *transition probability*
- Each emission is associated with an *emission probability*
- Two conditions:
  - All outgoing transition arcs from a state must sum to 1
  - All emission arcs from a state must sum to 1

2/20/08

6


## Probabilistic FSMs



$$\sum_x P(q \rightarrow x) = P(q \rightarrow r) + P(q \rightarrow q) = 1.0$$

$$\sum_x P(\text{emit}(q, x)) = P(\text{emit}(q, a)) + P(\text{emit}(q, b)) = 1.0$$

## Probabilistic FSMs



# Hidden Markov Models

- There are  $n$  states  $s_1, \dots, s_i, \dots, s_n$
- The emissions are observed (input data)
- Observation sequence  $\mathbf{O}=(o_1, \dots, o_i, \dots, o_T)$
- The states are not directly observed (hidden)
- Data does not directly tell us which state  $X_t$  is linked with observation  $o_t$


$$X_t \in \{s_1, \dots, s_n\}$$

2/20/08

9

## Markov Chains vs. HMMs

- For observation sequence *babaa*  
*i.e.*:  $o_1=b, o_2=a, \dots, o_5=a$
- Compute  $P(babaa)$  using a bigram model  
 $P(b)*P(a|b)*P(b|a)*P(a|b)*P(a|a)$
- Equivalent Markov chain:




2/20/08

10


## Markov Chains vs. HMMs

- For observation sequence  $babaa$   
i.e.:  $o_1=b, o_2=a, \dots, o_5=a$
- Compute  $P(babaa)$  using a trigram model  
 $P(ba)*P(b/ba)*P(a/lab)*P(alba)$
- Equivalent Markov chain:



## Markov Chains vs. HMMs

- For observation sequence  $babaa$   
i.e.:  $o_1=b, o_2=a, \dots, o_5=a$
- Compute  $P(babaa)$  using a trigram model  
 $P(ba)*P(b/ba)*P(a/lab)*P(alba)$
- Equivalent Markov chain:



## Markov Chains vs. HMMs

- Given an observation sequence  
 $\mathbf{O}=(o_1, \dots, o_p, \dots, o_T)$
- An  $n$ th order Markov Chain or  $n$ -gram model computes the probability  
 $P(o_1, \dots, o_p, \dots, o_T)$
- An HMM computes the probability  
 $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$  where the state sequence is *hidden*

2/20/08

13

## Properties of HMMs

- Markov assumption

$$P(X_t = s_i \mid \dots, X_{t-1} = s_j)$$

- Stationary distribution

$$P(X_t = s_i \mid X_{t-1} = s_j) = P(X_{t+l} = s_i \mid X_{t+l-1} = s_j)$$

2/20/08

14

# HMM Algorithms

- HMM as language model: compute probability of given observation sequence
- HMM as parser: compute the best sequence of states for a given observation sequence
- HMM as learner: given a set of observation sequences, learn its distribution, i.e. learn the transition and emission probabilities

2/20/08

15

# HMM Algorithms

- HMM as language model: compute probability of given observation sequence
- Compute  $P(o_1, \dots, o_T)$  from the probability  $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$

$$= \prod_{t=1}^T P(X_{t+1} = s_j \mid X_t = s_i) \times P(o_t = k \mid X_{t+1} = s_j)$$
$$P(o_1, \dots, o_T) = \sum_{X_1, \dots, X_{T+1}} P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$$

2/20/08

16



# HMM Algorithms


- HMM as parser: compute the best sequence of states for a given observation sequence
  - Compute best path  $X_1, \dots, X_{T+1}$  from the probability  $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$
- Best state sequence  $X^*_1, \dots, X^*_{T+1}$

$$= \operatorname{argmax}_{X_1, \dots, X_{T+1}} P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$$

2/20/08

17

## Best Path (Viterbi) Algorithm



- Key Idea 1: storing just the best path doesn't work
- Key Idea 2: store the best path upto *each* state

2/20/08

18

## Viterbi Algorithm

```
function viterbi (edges, input, obs): returns best path
edges = transition probability
input = emission probability
T = length of obs, the observation sequence
num-states = number of states in the HMM
Create a path-matrix: viterbi[num-states+1, T+1] # init to all 0s
for each state s: viterbi[s, 0] =  $\pi[s]$ 
for each time step t from 0 to T:
    for each state s from 0 to num-states:
        for each s' where edges[s,s'] is a transition probability:
            new-score = viterbi[s,t] * edges[s,s'] * input[s',obs[t]]
            if (viterbi[s',t+1] == 0) or (new-score > viterbi[s', t+1]):
                viterbi[s', t+1] = new-score
                back-pointer[s',t+1] = s
```

2/20/08

19

## Viterbi Algorithm

```
# finding the best path
best-final-score = best-final-state = 0
for each state s from 0 to num-states:
    if (viterbi[s,T+1] > best-final-score):
        best-final-state = s
        best-final-score = viterbi[s,T+1]
# start with the last state in the sequence
x = best-final-state
state-sequence.push(x)
for t from T+1 downto 0:
    state-sequence.push(back-pointer[x,t])
    x = back-pointer[x,t]
return state-sequence
```

2/20/08

20

## Forward-Backward Algorithm

- Algorithm that finds the transition and emission probabilities using training data that *does not have* hidden states provided
- Set the probabilities (for all parameters in the HMM) so that the training data T is assigned highest P(T) value (or lowest H(T), entropy value)
- This is called the maximum likelihood value over all possible hidden state sequences for the training data
- Exploits the fact that some transitions and resulting observations will occur more frequently than others in the training data

2/20/08

21

## Forward-backward Algorithm

- Consider input  $o_1, \dots, o_p, \dots, o_T$  where each  $o_t$  is from a set of symbols  $V = \{1, \dots, k, \dots, K\}$
- Let  $\pi_i$  be the probability of state  $i$  being a start state (for simplicity,  $\pi_i$  is not discussed further)
- Let  $a_{i,j}$  be the transition probability:  

$$P(X_{t+1} = s_j \mid X_t = s_i) \quad |S|^2 \text{ distinct } a_{i,j} \text{ values}$$
- Let  $b_{j,k}$  be the emission probability:  

$$P(o_t = k \mid X_{t+1} = s_j) \quad |S| \times |V| \text{ distinct } b_{j,k} \text{ values}$$
- Probability of going from state  $s_i$  to state  $s_j$  while observing input  $o_t$  is simply  $a_{i,j} \times b_{j,k}$

2/20/08

22

## Forward-backward Algorithm

- The algorithm starts with an initial setting for the probabilities in  $a$  and  $b$
- We are provided with training data which consists of observation sequence(s):  $o_1, \dots, o_t, \dots, o_T$
- The probability  $P(o_1, \dots, o_T)$  depends on the values in  $a$  and  $b$
- For given observation sequence(s), different transitions/emissions will be visited with different frequencies

2/20/08

23

## Forward-backward Algorithm

- For every path through the HMM, we count how many transitions occurred from state  $i$  to state  $j$  on observation  $o_t$
- Then (loosely speaking) we reward those transitions (and emissions) which have high *expected* frequency and penalize the competing transitions
- Expected frequency means we multiply the frequency with the current probability (taken from  $a$  and  $b$ )

2/20/08

24

## Forward-backward Algorithm

- $P(o_1, \dots, o_T)$  is the expected frequency of visiting all transitions and so the new frequency is the expected occurrence of a transition divided by  $P(o_1, \dots, o_T)$
- This gives us new values for all probabilities:  $a'$  and  $b'$  and we set  $a$  and  $b$  to these new values
- Compute  $P(o_1, \dots, o_T)$ . If the value is unchanged from before iteration then stop (convergence)
- Otherwise iterate (the entire procedure) with new values for  $a$  and  $b$

2/20/08

25


## Forward-backward Algorithm

- How to compute expected frequency over all paths efficiently (*reuse dynamic programming idea from Viterbi algorithm*)
- For input  $o_1, \dots, o_t, \dots, o_T$  where  $o_t \in V = \{1, \dots, k, \dots, K\}$
- For every path from a start state to state  $i$  we can compute the probability of observing  $o_1, \dots, o_{t-1}$
- Let  $\alpha_i(t)$  be the sum of all these probabilities
- For every path from state  $j$  to a final state we can compute the probability of observing  $o_{t+1}, \dots, o_T$
- Let  $\beta_j(t+1)$  be the sum of all these probabilities

2/20/08

26

## Forward-Backward Algorithm




$$\alpha_i(t) = \sum_{k=1}^N a_{k,i} \cdot b_{i,o_{t-1}} \cdot \alpha_k(t-1)$$

$$\beta_j(t+1) = \sum_{m=1}^N a_{j,m} \cdot b_{m,o_{t+1}} \cdot \beta_m(t+2)$$

$$P(o_1, \dots, o_T) = \sum_{i=1}^N \alpha_i(T+1) = \sum_{i=1}^N \pi_i \cdot \beta_i(1)$$

27

## Forward-Backward Algorithm



$$\hat{f}(i, j, o_t) = \frac{\alpha_i(t) \cdot a_{i,j} \cdot b_{j,o_t} \cdot \beta_j(t+1)}{P(o_1, \dots, o_T)} \quad \hat{f}(i, j) = \sum_{t=1}^T \hat{f}(i, j, o_t)$$

$$a'_{i,j} = \frac{\hat{f}(i, j)}{\sum_{j=1}^N \hat{f}(i, j)} \quad b'_{j,k} = \frac{\sum_{i=1}^N \sum_{t=1}^T \hat{f}(i, j, o_t = k)}{\sum_{i=1}^N \sum_{j=1}^N \hat{f}(i, j)}$$

2/20/08

28

# Forward-Backward Algorithm

- Each iteration provides new values for all the *parameters*
- But are the new parameters any better? How can we tell?
- Compute probability of the training data
- For HMMs, Baum 1977 shows that the probability will always be non-decreasing (later generalized to the more general EM algorithm)
- Same as cross-entropy is non-increasing

$$KL(\mu_{i+1} \parallel D) \leq KL(\mu_i \parallel D)$$