

Yet Another Introduction to Hidden Markov Models

Anoop Sarkar


School of Computing Science, Simon Fraser University

Burnaby, BC V5A 1S6, Canada

anoop@cs.sfu.ca

DOCTOR FUN

10 Feb 2000



Copyright © 2000 David Farley, d-farley@metabunc.edu
<http://metabunc.edu/Dave/drfun.html>
This cartoon is made available on the Internet for personal viewing only. Opinions expressed herein are solely those of the author.

Don Knuth finally sells out.

Abstract

Yet another introduction to Baum-Welch re-estimation algorithm for learning parameters of a Hidden Markov Model (HMM). The presentation is from an NLP perspective and the focus is on deriving efficient pseudo-code for a fast implementation of the Viterbi, forward and backward algorithms that are used in Baum-Welch re-estimation. We consider only discrete states and discrete observations. In fact, to make the presentation NLP friendly we assume the task will be to train from word sequences with tagging sequences as hidden data, where in the decoding step each word will be eventually tagged with its appropriate tag. The tags are discrete, but can refer to part of speech tags, chunking tags, tags that indicate word boundaries, and other NLP applications that can be represented as a sequence learning problem.

1 Notation

We refer to states or tags as $t \in \{1, \dots, N\}$ and observations or words as $w \in \{1, \dots, M\}$. In place of the tag index or word index we sometimes use an actual tag $t = \text{DT}$ or word $w = \text{the}$ in some examples. A special word $w = 0$ and tag $t = 0$ is used to indicate the begin and end of a sequence. The sequence of states and associated observations, or tags and associated words is $w_0, t_0, \dots, w_{T-1}, t_{T-1}$, a sequence of length T . We first consider a single sequence, later generalizing to multiple sequences of varying length.

2 Defining a Hidden Markov Model

A Hidden Markov Model (HMM) θ is the triple $\langle p_s, p_{tt}, p_{tw} \rangle$, where

1. $p_s(t_0)$ is the probability that we start with some tag t as t_0 ,
2. $p_{tt}(t_i | t_{i-1})$ is the transition probability from t_{i-1} to t_i , and
3. $p_{tw}(w_i | t_i)$ is the probability of generating w_i from t_i .

We set $p_s(0) = 1.0$ and $p_{tw}(0 | 0) = 1.0$. As a result we can now ignore p_s and our HMM is represented by the tuple $\langle p_{tt}, p_{tw} \rangle$.

3 The Forward Algorithm

Forward step (α):

$$\begin{aligned} \Pr(w_0 = 0 | t_0 = 0) &= 1.0 \\ \Pr(w_0, \dots, w_i, t_i) &= \sum_{t_{i-1}} \Pr(w_0, \dots, w_{i-1}, t_{i-1}) \cdot p_{tt}(t_i | t_{i-1}) \cdot p_{tw}(w_i | t_i) \\ \Pr(w_0, \dots, w_{T-1}) &= \sum_{t_{T-1}} \Pr(w_0, \dots, w_{T-1}, t_{T-1}) \\ &= \Pr(w_0, \dots, w_{T-1}, t_{T-1} = 0) \end{aligned} \tag{1}$$

$$\begin{aligned} \alpha(0, 0) &= 1.0 \\ \alpha(t_i, i) &= \sum_{t_{i-1}} \alpha(t_{i-1}, i-1) \cdot p_{tt}(t_i | t_{i-1}) \cdot p_{tw}(w_i | t_i), \quad 1 \leq i \leq T-1 \\ \Pr(w_0, \dots, w_{T-1}) &= \sum_t \alpha(t, T-1) \\ &= \alpha(0, T-1) \end{aligned} \tag{2}$$

Algorithm 1 α : implements the forward algorithm

Require: $W = w_0, \dots, w_{T-1}$

Require: $\text{tag-dict}: \{1, \dots, M\} \rightarrow \mathcal{P}(\{1, \dots, N\})$

```

1:  $\alpha(0, 0) := 1.0$ 
2: for  $(1 \leq i \leq T - 1)$  do
3:   for  $(t_i \in \text{tag-dict}(w_i))$  do
4:     for  $(t_{i-1} \in \text{tag-dict}(w_{i-1}))$  do
5:        $\alpha(t_i, i) := \alpha(t_i, i) + \alpha(t_{i-1}, i - 1) \cdot p_{\text{tt}}(t_i \mid t_{i-1}) \cdot p_{\text{tw}}(w_i \mid t_i)$ 
6:     end for
7:   end for
8: end for
9:  $s := \alpha(0, T - 1)$ 
10: return  $s$ 

```

Algorithm 2 α : implements the forward algorithm, 1st optimization

Require: $W = w_0, \dots, w_{T-1}$

Require: $\text{tag-dict}: \{1, \dots, M\} \rightarrow \mathcal{P}(\{1, \dots, N\})$

```

1:  $\alpha(0, 0) := 1.0$ 
2: for  $(1 \leq i \leq T - 1)$  do
3:   for  $(t_i \in \text{tag-dict}(w_i))$  do
4:     for  $(t_{i-1} \in \text{tag-dict}(w_{i-1}))$  do
5:        $\alpha(t_i, i) := \alpha(t_i, i) + \alpha(t_{i-1}, i - 1) \cdot p_{\text{tt}}(t_i \mid t_{i-1})$ 
6:     end for
7:      $\alpha(t_i, i) := \alpha(t_i, i) \cdot p_{\text{tw}}(w_i \mid t_i)$ 
8:   end for
9: end for
10:  $s := \alpha(0, T - 1)$ 
11: return  $s$ 

```

4 The Backward Algorithm

Backward step (β):

$$\begin{aligned}
\Pr(w_{T-1} = 0 \mid t_{T-2}) &= 1.0 \quad \text{for all } t_{T-2} \\
\Pr(w_i, \dots, w_{T-1} \mid t_{i-1}) &= \sum_{t_i} \Pr(w_{i+1}, \dots, w_{T-1} \mid t_i) \cdot p_{\text{tt}}(t_i \mid t_{i-1}) \cdot p_{\text{tw}}(w_i \mid t_i) \\
\Pr(w_0, \dots, w_{T-1}) &= \sum_{t_0} \Pr(w_1, \dots, w_{T-1} \mid t_0) \\
&= \Pr(w_1, \dots, w_{T-1} \mid 0)
\end{aligned} \tag{3}$$

$$\begin{aligned}
\beta(0, T - 1) &= 1.0 \\
\beta(t_{i-1}, i - 1) &= \sum_{t_i} \beta(t_i, i) \cdot p_{\text{tt}}(t_i \mid t_{i-1}) \cdot p_{\text{tw}}(w_i \mid t_i), \quad T - 1 \geq i > 0 \\
\Pr(w_0, \dots, w_{T-1}) &= \sum_t \beta(t, 0) \\
&= \beta(0, 0)
\end{aligned} \tag{4}$$

Algorithm 3 β : implements the backward algorithm

Require: $W = w_0, \dots, w_{T-1}$

Require: $\text{tag-dict}: \{1, \dots, M\} \rightarrow \mathcal{P}(\{1, \dots, N\})$

```

1:  $\beta(0, T-1) := 1.0$ 
2: for  $(T-1 \geq i \geq 0)$  do
3:   for  $(t_i \in \text{tag-dict}(w_i))$  do
4:     for  $(t_{i-1} \in \text{tag-dict}(w_{i-1}))$  do
5:        $\beta(t_{i-1}, i-1) := \beta(t_{i-1}, i-1) + p_{\text{tt}}(t_i | t_{i-1}) \cdot p_{\text{tw}}(w_i | t_i) \cdot \beta(t_i, i)$ 
6:     end for
7:   end for
8: end for

```

5 Baum-Welch Re-estimation

Computing γ :

$$\begin{aligned}
\gamma(i, t_i) &= \Pr(t_i | w_0, \dots, w_{T-1}) \text{ for } t_i \in \{1, \dots, N\} \\
&= \frac{\Pr(w_0, \dots, w_i, t_i) \cdot \Pr(w_{i+1}, \dots, w_{T-1} | t_i)}{\sum_{s_i} \Pr(w_0, \dots, w_i, s_i) \cdot \Pr(w_{i+1}, \dots, w_{T-1} | s_i)} \\
&= \frac{\Pr(w_0, \dots, w_i, t_i) \cdot \Pr(w_{i+1}, \dots, w_{T-1} | t_i)}{\Pr(w_0, \dots, w_{T-1})} \\
&= \frac{\alpha(t_i, i) \cdot \beta(t_i, i)}{\alpha(0, T-1)} \tag{5}
\end{aligned}$$

Computing ξ :

$$\begin{aligned}
\xi(i, t_i, t_{i+1}) &= \Pr(t_i, t_{i+1} | w_0, \dots, w_{T-1}) \text{ for } t_i, t_{i+1} \in \{1, \dots, N\} \\
&= \frac{\Pr(w_0, \dots, w_i, t_i) \cdot p_{\text{tt}}(t_{i+1} | t_i) \cdot p_{\text{tw}}(w_{i+1} | t_{i+1}) \cdot \Pr(w_{i+2}, \dots, w_{T-1} | t_{i+1})}{\sum_{s_i, s_{i+1}} \Pr(w_0, \dots, w_i, s_i) \cdot p_{\text{tt}}(s_{i+1} | s_i) \cdot p_{\text{tw}}(w_{i+1} | s_{i+1}) \cdot \Pr(w_{i+2}, \dots, w_{T-1} | s_{i+1})} \\
&= \frac{\alpha(t_i, i) \cdot p_{\text{tt}}(t_{i+1} | t_i) \cdot p_{\text{tw}}(w_{i+1} | t_{i+1}) \cdot \beta(t_{i+1}, i+1)}{\sum_{s_i, s_{i+1}} \alpha(s_i, i) \cdot p_{\text{tt}}(s_{i+1} | s_i) \cdot p_{\text{tw}}(w_{i+1} | s_{i+1}) \cdot \beta(s_{i+1}, i+1)} \\
&= \frac{\alpha(t_i, i) \cdot p_{\text{tt}}(t_{i+1} | t_i) \cdot p_{\text{tw}}(w_{i+1} | t_{i+1}) \cdot \beta(t_{i+1}, i+1)}{\Pr(w_0, \dots, w_{T-1})} \\
&= \frac{\alpha(t_i, i) \cdot p_{\text{tt}}(t_{i+1} | t_i) \cdot p_{\text{tw}}(w_{i+1} | t_{i+1}) \cdot \beta(t_{i+1}, i+1)}{\alpha(0, T-1)} \tag{6}
\end{aligned}$$

To save space in computing ξ , we compute ξ' :

$$\begin{aligned}
\xi'(t_i, t_{i+1}) &= \sum_{i=0}^{T-2} \frac{\alpha(t_i, i) \cdot p_{\text{tt}}(t_{i+1} | t_i) \cdot p_{\text{tw}}(w_{i+1} | t_{i+1}) \cdot \beta(t_{i+1}, i+1)}{\alpha(0, T-1)} \\
&\quad \text{for all } t_i, t_{i+1} \in \{1, \dots, N\} \tag{7}
\end{aligned}$$

and δ' :

$$\delta'(t, w) = \sum_{i=0}^{T-1} \gamma(i, t_i = t) \cdot I(w_i = w)$$

$$\begin{aligned} &\text{where } I(\text{true}) = 1 \text{ and } I(\text{false}) = 0. \\ &\text{for all } t \in \{1, \dots, N\}, w \in \{1, \dots, M\} \end{aligned} \quad (8)$$

The new estimates for p_{tt} and p_{tw} are:

$$p_{\text{tt}}(t_{i+1} \mid t_i) = \frac{\xi'(t_i, t_{i+1})}{\sum_{s_i} \xi'(t_i, s_i)} \quad (9)$$

$$p_{\text{tw}}(w \mid t_i) = \frac{\delta'(t_i, w)}{\sum_{w'} \delta'(t_i, w')} \quad (10)$$

Algorithm 4 Forward-Backward:

1 iteration of iterative algorithm to get new values for p_{tt} and p_{tw} ,
computes β at the same time

Require: $W = w_0, \dots, w_{T-1}$

Require: $\text{tag-dict}: \{1, \dots, M\} \rightarrow \mathcal{P}(\{1, \dots, N\})$

```

1:  $s := \alpha(W, \text{tag-dict})$ 
2:  $\beta(0, T-1) := 1.0$ 
3: for  $(T-1 \geq i \geq 0)$  do
4:   for  $(t_i \in \text{tag-dict}(w_i))$  do
5:      $\delta'(t_i, w_i) := \delta'(t_i, w_i) + \frac{1}{s} \cdot \alpha(i, t_i) \cdot \beta(i, t_i)$ 
6:      $\delta''(t_i) := \delta''(t_i) + \delta'(t_i, w_i)$ 
7:     for  $(t_{i-1} \in \text{tag-dict}(w_{i-1}))$  do
8:        $p := p_{\text{tt}}(t_i \mid t_{i-1}) \cdot p_{\text{tw}}(w_i \mid t_i)$ 
9:        $\beta(t_{i-1}, i-1) := \beta(t_{i-1}, i-1) + p \cdot \beta(t_i, i)$ 
10:       $\xi'(t_{i-1}, t_i) := \xi'(t_{i-1}, t_i) + \frac{1}{s} \cdot \alpha(t_{i-1}, i-1) \cdot p \cdot \beta(t_i, i)$ 
11:       $\xi''(t_{i-1}) := \xi''(t_{i-1}) + \xi'(t_{i-1}, t_i)$ 
12:    end for
13:  end for
14: end for
15: for  $(t_i \in \{1, \dots, N\})$  do
16:   for  $(w \in \{1, \dots, M\})$  do
17:      $p_{\text{tw}}(w \mid t_i) := \frac{\delta'(t_i, w)}{\delta''(t_i)}$ 
18:   end for
19:   for  $(t_{i+1} \in \{1, \dots, N\})$  do
20:      $p_{\text{tt}}(t_{i+1} \mid t_i) := \frac{\xi'(t_i, t_{i+1})}{\xi''(t_i)}$ 
21:   end for
22: end for
```

6 Conclusion

Did you understand it?

References

- Jason Eisner. 2004. *Tagging with a Hidden Markov Model. Lecture Notes from Johns Hopkins course 600.465 — Introduction to NLP*
- John Langford. 2000. *Optimizing Hidden Markov Model Learning*. unpublished manuscript.

Algorithm 5 Forward-Backward: 1st optimization

Require: $W = w_0, \dots, w_{T-1}$

Require: $\text{tag-dict}: \{1, \dots, M\} \rightarrow \mathcal{P}(\{1, \dots, N\})$

```
1:  $s := \alpha(W, \text{tag-dict})$ 
2:  $\beta(0, T-1) := 1.0$ 
3: for  $(T-1 \geq i \geq 0)$  do
4:   for  $(t_i \in \text{tag-dict}(w_i))$  do
5:      $\delta'(t_i, w_i) := \delta'(t_i, w_i) + \frac{1}{s} \cdot \alpha(i, t_i) \cdot \beta(i, t_i)$ 
6:      $\delta''(t_i) := \delta''(t_i) + \delta'(t_i, w_i)$ 
7:      $m := p_{\text{tw}}(w_i \mid t_i) \cdot \beta(t_i, i)$ 
8:     for  $(t_{i-1} \in \text{tag-dict}(w_{i-1}))$  do
9:        $n := p_{\text{tt}}(t_i \mid t_{i-1}) \cdot m$ 
10:       $\beta(t_{i-1}, i-1) := \beta(t_{i-1}, i-1) + n$ 
11:       $\xi'(t_{i-1}, t_i) := \xi'(t_{i-1}, t_i) + \frac{1}{s} \cdot \alpha(t_{i-1}, i-1) \cdot n$ 
12:       $\xi''(t_{i-1}) := \xi''(t_{i-1}) + \xi'(t_{i-1}, t_i)$ 
13:    end for
14:  end for
15: end for
16: for  $(t_i \in \{1, \dots, N\})$  do
17:   for  $(w \in \{1, \dots, M\})$  do
18:      $p_{\text{tw}}(w \mid t_i) := \frac{\delta'(t_i, w)}{\delta''(t_i)}$ 
19:   end for
20:   for  $(t_{i+1} \in \{1, \dots, N\})$  do
21:      $p_{\text{tt}}(t_{i+1} \mid t_i) := \frac{\xi'(t_i, t_{i+1})}{\xi''(t_i)}$ 
22:   end for
23: end for
```

Algorithm 6 Forward-Backward: multiple training sentences

Require: $S = S_0, \dots, S_{K-1}$, a list of training sentences

Require: $tag\text{-}dict: \{1, \dots, M\} \rightarrow \mathcal{P}(\{1, \dots, N\})$

```
1: for ( $0 \leq k \leq K - 1$ ) do
2:    $S_k = w_0, \dots, w_{T-1}$ 
3:    $s := \alpha(S_k, tag\text{-}dict)$ 
4:    $L := L \cdot s$ 
5:    $\beta(0, T - 1) := 1.0$ 
6:   for ( $T - 1 \geq i \geq 0$ ) do
7:     for ( $t_i \in tag\text{-}dict(w_i)$ ) do
8:        $\delta'(t_i, w_i) := \delta'(t_i, w_i) + \frac{1}{s} \cdot \alpha(i, t_i) \cdot \beta(i, t_i)$ 
9:        $\delta''(t_i) := \delta''(t_i) + \delta'(t_i, w_i)$ 
10:       $m := p_{\text{tw}}(w_i \mid t_i) \cdot \beta(t_i, i)$ 
11:      for ( $t_{i-1} \in tag\text{-}dict(w_{i-1})$ ) do
12:         $n := p_{\text{tt}}(t_i \mid t_{i-1}) \cdot m$ 
13:         $\beta(t_{i-1}, i - 1) := \beta(t_{i-1}, i - 1) + n$ 
14:         $\xi'(t_{i-1}, t_i) := \xi'(t_{i-1}, t_i) + \frac{1}{s} \cdot \alpha(t_{i-1}, i - 1) \cdot n$ 
15:         $\xi''(t_{i-1}) := \xi''(t_{i-1}) + \xi'(t_{i-1}, t_i)$ 
16:      end for
17:    end for
18:  end for
19: end for
20: for ( $t_i \in \{1, \dots, N\}$ ) do
21:   for ( $w \in \{1, \dots, M\}$ ) do
22:      $p_{\text{tw}}(w \mid t_i) := \frac{\delta'(t_i, w)}{\delta''(t_i)}$ 
23:   end for
24:   for ( $t_{i+1} \in \{1, \dots, N\}$ ) do
25:      $p_{\text{tt}}(t_{i+1} \mid t_i) := \frac{\xi'(t_i, t_{i+1})}{\xi''(t_i)}$ 
26:   end for
27: end for
28: return  $L$ , the likelihood of the training set  $S$ 
```

Algorithm 7 Baum-Welch Re-estimation: using only unlabelled data

Require: $S = S_0, \dots, S_{K-1}$, a list of unlabelled training sentences

Require: $threshold$: small positive number, e.g. 10^{-5} , change in likelihood for stopping condition

```
1: for ( $w \in S$ ) do
2:    $tag\text{-}dict(w) := \{1, \dots, N\}$ 
3: end for
4:  $L := 0$ 
5: repeat
6:    $L' := L$ 
7:    $L := Forward\text{-}Backward(S, tag\text{-}dict)$ 
8:   assert ( $L - L' > 0$ )
9: until ( $L - L' > threshold$ )
10: return  $p_{\text{tt}}, p_{\text{tw}}$ 
```
