

CMPT-413

Computational Linguistics

Anoop Sarkar
<http://www.cs.sfu.ca/~anoop>

March 5, 2012

1 / 27

Prepositional Phrases

- ▶ noun attach: *I bought the shirt with pockets*
- ▶ verb attach: *I washed the shirt with soap*
- ▶ As in the case of other attachment decisions in parsing: it depends on the meaning of the entire sentence – needs world knowledge, etc.
- ▶ Maybe there is a simpler solution: we can attempt to solve it using heuristics or associations between words

2 / 27

Structure Based Ambiguity Resolution

- ▶ Right association: a constituent (NP or PP) tends to attach to another constituent immediately to its right (Kimball 1973)
- ▶ Minimal attachment: a constituent tends to attach to an existing non-terminal using the fewest additional syntactic nodes (Frazier 1978)
- ▶ These two principles make opposite predictions for prepositional phrase attachment
- ▶ Consider the grammar:

$$VP \rightarrow V NP PP \quad (1)$$

$$NP \rightarrow NP PP \quad (2)$$

for input: *I* [_{VP} *saw* [_{NP} *the man* ... [_{PP} *with the telescope*],
RA predicts that the PP attaches to the NP, i.e. use rule (2),
and MA predicts V attachment, i.e. use rule (1)

3 / 27

Structure Based Ambiguity Resolution

- ▶ Garden-paths look structural:
The emergency crews hate most is domestic violence
- ▶ Neither MA or RA account for more than 55% of the cases in real text
- ▶ Psycholinguistic experiments using eyetracking show that humans resolve ambiguities as soon as possible in the left to right sequence using the words to disambiguate
- ▶ Garden-paths are caused by a combination of lexical and structural effects:
The flowers delivered for the patient arrived

4 / 27

Ambiguity Resolution: Prepositional Phrases in English

► Learning Prepositional Phrase Attachment: Annotated Data

v	n1	p	n2	Attachment
join	board	as	director	V
is	chairman	of	N.V.	N
using	crocidolite	in	filters	V
bring	attention	to	problem	V
is	asbestos	in	products	N
making	paper	for	filters	N
including	three	with	cancer	N
⋮	⋮	⋮	⋮	⋮

5 / 27

Prepositional Phrase Attachment

Method	Accuracy
Always noun attachment	59.0
Most likely for each preposition	72.2
Average Human (4 head words only)	88.2
Average Human (whole sentence)	93.2

6 / 27

Back-off Smoothing

- ▶ Let 1 represent noun attachment.
- ▶ We want to compute probability of noun attachment:
 $p(1 \mid v, n1, p, n2)$.
- ▶ Probability of verb attachment is $1 - p(1 \mid v, n1, p, n2)$.

7 / 27

Back-off Smoothing

1. If $f(v, n1, p, n2) > 0$ and $\hat{p} \neq 0.5$

$$\hat{p}(1 \mid v, n1, p, n2) = \frac{f(1, v, n1, p, n2)}{f(v, n1, p, n2)}$$

2. Else if $f(v, n1, p) + f(v, p, n2) + f(n1, p, n2) > 0$
and $\hat{p} \neq 0.5$

$$\hat{p}(1 \mid v, n1, p, n2) = \frac{f(1, v, n1, p) + f(1, v, p, n2) + f(1, n1, p, n2)}{f(v, n1, p) + f(v, p, n2) + f(n1, p, n2)}$$

3. Else if $f(v, p) + f(n1, p) + f(p, n2) > 0$

$$\hat{p}(1 \mid v, n1, p, n2) = \frac{f(1, v, p) + f(1, n1, p) + f(1, p, n2)}{f(v, p) + f(n1, p) + f(p, n2)}$$

4. Else if $f(p) > 0$

$$\hat{p}(1 \mid v, n1, p, n2) = \frac{f(1, p)}{f(p)}$$

5. Else $\hat{p}(1 \mid v, n1, p, n2) = 1.0$

8 / 27

Prepositional Phrase Attachment: (Collins and Brooks 1995)

- ▶ **Results:** 84.5% accuracy
with the use of some limited word classes for dates, numbers, etc.
- ▶ Using complex word classes taken from WordNet (which we shall be looking at later in this course) increases accuracy to 88% (Stetina and Nagao 1998)
- ▶ We can improve on parsing performance with Probabilistic CFGs by using the insights taken from PP attachment.
- ▶ Modify the PCFG model to be sensitive to words and other context-sensitive features of the input.
- ▶ And generalizing to other kinds of attachment problems, like coordination or deciding which constituent is an argument of a verb.

9 / 27

Some other studies

- ▶ **Toutanova, Manning, and Ng, 2004:**
use sophisticated smoothing model for PP attachment
86.18% with words & stems; with word classes: 87.54%
- ▶ **Merlo, Crocker and Berthouzoz, 1997:**
test on multiple PPs, generalize disambiguation of 1 PP to 2-3 PPs
14 structures possible for 3PPs assuming a single verb: all 14 are attested in the Treebank
same model as CB95; but generalized to dealing with upto 3PPs
1PP: 84.3% 2PP: 69.6% 3PP: 43.6%
Note that this is still not the real problem faced in parsing natural language

10 / 27

Probability Models

- ▶ $p(x, y)$: x = input, y = labels
- ▶ Pick best prob distribution $p(x, y)$ to fit the data
- ▶ Max likelihood of the data *according to the prob model* equivalent to minimizing entropy

11 / 27

Probability Models

- ▶ Max likelihood of the data *according to the prob model*
- ▶ Equivalent to picking best parameter values θ such that the data gets highest likelihood:

$$\max_{\theta} p(\theta \mid \text{data}) = \max_{\theta} p(\theta) \cdot p(\text{data} \mid \theta)$$

12 / 27

Advantages of probability models

- ▶ parameters can be estimated automatically, while scores have to be twiddled by hand
- ▶ parameters can be estimated from supervised or unsupervised data
- ▶ probabilities can be used to quantify confidence in a particular state and used to compare against other probabilities in a strictly comparable setting
- ▶ modularity: $p(\text{semantics}) \cdot p(\text{syntax} \mid \text{semantics}) \cdot p(\text{morphology} \mid \text{syntax}) \cdot p(\text{phonology} \mid \text{morphology}) \cdot p(\text{sounds} \mid \text{phonology})$

13 / 27

Naive Bayes Classifier

- ▶ \mathbf{x} is the input that can be represented as d independent features f_j , $1 \leq j \leq d$
- ▶ y is the output classification
- ▶ $P(y \mid \mathbf{x}) = \frac{P(y) \cdot P(\mathbf{x} \mid y)}{P(\mathbf{x})}$
- ▶ $P(\mathbf{x} \mid y) = \prod_{j=1}^d P(f_j \mid y)$
- ▶ $P(y \mid \mathbf{x}) = P(y) \cdot \prod_{j=1}^d P(f_j \mid y)$

14 / 27

Using Naive Bayes for Document Classification

- ▶ Spam text: Learn how to make \$38.99 into a money making machine that pays ... \$7,000 / month !
- ▶ Distinguish spam text from regular email text
- ▶ Find useful features to make this distinction

15 / 27

Using Naive Bayes

- ▶ Useful features
 1. contains turn \$AMOUNT into
 2. contains \$AMOUNT
 3. contains Learn how to
 4. contains exclamation mark at end of sentence

16 / 27

Using Naive Bayes

- ▶ how many times do these features occur?
 1. contains: turn \$AMOUNT into
 - in spam text: 50
 - in normal email: 2
 - i.e. 25x more likely in spam
 2. contains: \$AMOUNT
 - in spam text: 90
 - in normal email: 10
 - i.e. 9x more likely in spam

17 / 27

Using Naive Bayes

- ▶ How likely is it for *both* features to occur at the same time in a spam message?
 1. contains: turn \$AMOUNT into
 2. contains: \$AMOUNT
- ▶ Assume we have a new feature, contains: turn \$AMOUNT into *and* \$AMOUNT
- ▶ The model predicts that the event that both features occur simultaneously has probability $\frac{140}{152} = 0.92$
- ▶ But Naive Bayes assumes that these features are independent and should occur with probability:
 $0.92 \cdot 0.9 = 0.864$

18 / 27

Using Naive Bayes

- ▶ Naive Bayes needs overlapping but independent features
- ▶ How can we use all of the features we want?
 1. contains turn \$AMOUNT into
 2. contains \$AMOUNT
 3. contains Learn how to
 4. contains exclamation mark at end of sentence
- ▶ how about giving each feature a weight w equal to its log probability: $w = \log p(f, y)$

19 / 27

Using Naive Bayes

- ▶ each feature gets a score equal to its log probability
- ▶ Assign scores to features:
 1. $w_1 = +1$ contains turn \$AMOUNT into
 2. $w_2 = +5$ contains \$AMOUNT
 3. $w_3 = +0.2$ contains Learn how to
 4. $w_4 = -2$ contains exclamation mark at end of sentence

20 / 27

Using Naive Bayes

- ▶ so add the scores and treat it like a log probability
- ▶ $\log p(\text{spam} \mid \text{feats}) = 4.2$
- ▶ but then, $p(\text{spam} \mid \text{feats}) = \exp(4.2) = 66.68$
- ▶ how do we compute keep arbitrary scores and still get probabilities?

21 / 27

Log linear model

- ▶ Let there be m features, $f_k(\mathbf{x}, y)$ for $k = 1, \dots, m$
- ▶ Define a parameter vector $\mathbf{w} \in \mathbb{R}^m$
- ▶ Each (\mathbf{x}, y) pair is mapped to score:

$$s(\mathbf{x}, y) = \sum_k w_k \cdot f_k(\mathbf{x}, y)$$

- ▶ Using inner product notation:

$$\begin{aligned}\mathbf{w} \cdot \mathbf{f}(\mathbf{x}, y) &= \sum_k w_k \cdot f_k(\mathbf{x}, y) \\ s(\mathbf{x}, y) &= \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, y)\end{aligned}$$

- ▶ To get a probability from the score: Renormalize!

$$\Pr(y \mid \mathbf{x}, \mathbf{w}) = \frac{\exp(s(\mathbf{x}, y))}{\sum_{y'} \exp(s(\mathbf{x}, y'))}$$

22 / 27

Log linear model

- The name ‘log-linear model’ comes from:

$$\log \Pr(y \mid \mathbf{x}, \mathbf{w}) = \underbrace{\mathbf{w} \cdot \mathbf{f}(\mathbf{x}, y)}_{\text{linear term}} - \underbrace{\log \sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(\mathbf{x}, y'))}_{\text{normalization term}}$$

- Once the weights are learned, we can perform predictions using these features.
- The goal: to find \mathbf{w} that maximizes the log likelihood $L(\mathbf{w})$ of the labeled training set containing (\mathbf{x}_i, y_i) for $i = 1 \dots n$

$$\begin{aligned} L(\mathbf{w}) &= \sum_i \log \Pr(y_i \mid \mathbf{x}_i, \mathbf{w}) \\ &= \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y_i) - \sum_i \log \sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y')) \end{aligned}$$

23 / 27

Log linear model

- Maximize:

$$L(\mathbf{w}) = \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y_i) - \sum_i \log \sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y'))$$

- Calculate gradient:

$$\begin{aligned} &\left. \frac{dL(\mathbf{w})}{d\mathbf{w}} \right|_{\mathbf{w}} \\ &= \sum_i \mathbf{f}(\mathbf{x}_i, y_i) - \sum_i \frac{1}{\sum_{y''} \exp(\mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y''))} \\ &\quad \sum_{y'} \mathbf{f}(\mathbf{x}_i, y') \cdot \exp(\mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y')) \\ &= \sum_i \mathbf{f}(\mathbf{x}_i, y_i) - \sum_i \sum_{y'} \mathbf{f}(\mathbf{x}_i, y') \frac{\exp(\mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y'))}{\sum_{y''} \exp(\mathbf{w} \cdot \mathbf{f}(\mathbf{x}_i, y''))} \\ &= \underbrace{\sum_i \mathbf{f}(\mathbf{x}_i, y_i)}_{\text{Observed counts}} - \underbrace{\sum_i \sum_{y'} \mathbf{f}(\mathbf{x}_i, y') \Pr(y' \mid \mathbf{x}_i, \mathbf{w})}_{\text{Expected counts}} \end{aligned}$$

24 / 27

Log linear model

- ▶ Init: $\mathbf{w}^{(0)} = \mathbf{0}$
- ▶ $t \leftarrow 0$
- ▶ Iterate until convergence:
 - ▶ Calculate: $\Delta = \left. \frac{dL(\mathbf{w})}{d\mathbf{w}} \right|_{\mathbf{w}=\mathbf{w}^{(t)}}$
 - ▶ Find $\beta^* = \operatorname{argmax}_{\beta} L(\mathbf{w}^{(t)} + \beta\Delta)$
 - ▶ Set $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} + \beta^* \Delta$

25 / 27

Learning the weights: \mathbf{w} : Generalized Iterative Scaling

$$f^{\#} = \max_{x,y} \sum_{j=1}^k f_j(x, y)$$

For each iteration

expected[1 .. # of features] $\leftarrow 0$

For $i = 1$ to | training data |

For each feature f_j

expected[j] $+= f_j(x_i, y_i) \cdot P(y_i | x_i)$

For each feature f_j

observed[j] $= f_j(x, y) \cdot \frac{c(x,y)}{|\text{training data}|}$

For each feature f_j

$w_j \leftarrow w_j \cdot \sqrt{\frac{f^{\#} \cdot \text{observed[j]}}{\text{expected[j]}}}$

cf. Goodman, NIPS '01

26 / 27

Maximum Entropy

- ▶ The log-linear model has an interpretation as a *maximum entropy* model.
- ▶ For observed events, maximize likelihood. For unobserved events, from all consistent models pick the one with maximum entropy.
- ▶ The maximum entropy principle: related to Occam's razor and other similar justifications for scientific inquiry
- ▶ Make the minimum possible assumptions about unseen data
- ▶ Also: Laplace's *Principle of Insufficient Reason*: when one has no information to distinguish between the probability of two events, the best strategy is to consider them equally likely