

CMPT 413

Computational Linguistics

Anoop Sarkar

<http://www.cs.sfu.ca/~anoop>

3/19/12

1


Sequence Learning

- British Left Waffles on Falkland Islands
 - (N, N, V, P, N, N)
 - (N, V, N, P, N, N)
- Segmentation 中国十四个边境开放城市经济建设成就显著
 - (b, i, b, i, b, b, i, b, i, b, i, b, i, b, i, b, i)
 - 中国 十 四 个 边 境 开 放 城 市 经 济 建 设 成 就 显 著
 - China 's 14 open border cities marked economic achievements

3/19/12

2

Sequence Learning



3 states: N, V, P


Observation sequence: (o_1, \dots, o_6)

State sequence (6+1): $(Start, N, N, V, P, N, N)$

3/19/12

3

Finite State Machines




Mealy Machine

3/19/12

4

Finite State Machines



Moore Machine

3/19/12

5


Probabilistic FSMs

- Start at a state i with a *start state probability*: π_i
- Transition from state i to state j is associated with a *transition probability*: a_{ij}
- Emission of symbol o from state i is associated with an *emission probability*: $b_i(o)$
- Two conditions:
 - All outgoing transition arcs from a state must sum to 1
 - All symbol emissions from a state must sum to 1

3/19/12

6

Probabilistic FSMs




3/19/12

7

Probabilistic FSMs

Emission

$b_A(\text{killer}) = 0$
 $b_A(\text{crazy}) = 1$
 $b_A(\text{clown}) = 0$
 $b_A(\text{problem}) = 0$



Emission

$b_N(\text{killer}) = 1/3$
 $b_N(\text{crazy}) = 0$
 $b_N(\text{clown}) = 1/3$
 $b_N(\text{problem}) = 1/3$
 $\sum_{o \in V} b_i(o) = 1$

Start state

$\sum_i \pi_i = 1$
 $\pi_A = 1/2$
 $\pi_N = 1/2$

Transition

$a_{A,A} = 1/3$
 $a_{A,N} = 2/3$
 $a_{N,N} = 9/10$
 $a_{N,A} = 1/10$
 $\sum_j a_{i,j} = 1$

3/19/12

8

Hidden Markov Models


- There are n states $s_1, \dots, s_i, \dots, s_n$
- The emissions are observed (input data)
- Observation sequence $\mathbf{O}=(o_1, \dots, o_t, \dots, o_T)$
- The states are not directly observed (hidden)
- Data does not directly tell us which state X_t is linked with observation o_t
 $X_t \in \{s_1, \dots, s_n\}$

3/19/12

9

Markov Chains vs. HMMs

- For observation sequence *babaa*
i.e: $o_1=b, o_2=a, \dots, o_5=a$
- Compute $P(babaa)$ using a bigram model
 $P(b)*P(a|b)*P(b|a)*P(a|b)*P(a|a)$
- Equivalent Markov chain:




3/19/12

10


Markov Chains vs. HMMs

- For observation sequence $babaa$
i.e.: $o_1=b, o_2=a, \dots, o_5=a$
- Compute $P(babaa)$ using a trigram model
 $P(ba)*P(blba)*P(alab)*P(alba)$
- Equivalent Markov chain:



Markov Chains vs. HMMs

- For observation sequence $babaa$
i.e.: $o_1=b, o_2=a, \dots, o_5=a$
- Compute $P(babaa)$ using a trigram model
 $P(ba)*P(blba)*P(alab)*P(alba)$
- Equivalent Markov chain:



Markov Chains vs. HMMs

- Given an observation sequence
 $\mathbf{O}=(o_1, \dots, o_t, \dots, o_T)$
- An n th order Markov Chain or n -gram model computes the probability
 $P(o_1, \dots, o_t, \dots, o_T)$
- An HMM computes the probability
 $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$ where the state sequence is *hidden*

3/19/12

13

Properties of HMMs

- Markov assumption

$$P(X_t = s_i \mid \dots, X_{t-1} = s_j)$$

- Stationary distribution

$$P(X_t = s_i \mid X_{t-1} = s_j) = P(X_{t+l} = s_i \mid X_{t+l-1} = s_j)$$

3/19/12

14

HMM Algorithms

- HMM as language model: compute probability of given observation sequence
- HMM as parser: compute the best sequence of states for a given observation sequence
- HMM as learner: given a set of observation sequences, learn its distribution, i.e. learn the transition and emission probabilities

3/19/12

15

HMM Algorithms

- HMM as language model: compute probability of given observation sequence
- Compute $P(o_1, \dots, o_T)$ from the probability $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$
$$= \prod_{t=1}^T P(X_{t+1} = s_j \mid X_t = s_i) \times P(o_t = k \mid X_{t+1} = s_j)$$
$$P(o_1, \dots, o_T) = \sum_{X_1, \dots, X_{T+1}} P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$$

3/19/12

16

HMM Algorithms


- HMM as parser: compute the best sequence of states for a given observation sequence
 - Compute best path X_1, \dots, X_{T+1} from the probability $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$
- Best state sequence X_1^*, \dots, X_{T+1}^*

$$= \operatorname{argmax}_{X_1, \dots, X_{T+1}} P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$$

3/19/12

17

Best Path (Viterbi) Algorithm



- Key Idea 1: storing just the best path doesn't work
- Key Idea 2: store the best path upto *each* state

3/19/12

18

Viterbi Algorithm

```
function viterbi (edges, input, obs): returns best path
edges = transition probability
input = emission probability
T = length of obs, the observation sequence
num-states = number of states in the HMM
Create a path-matrix: viterbi[num-states+1, T+1] # init to all 0s
for each state s: viterbi[s, 0] =  $\pi[s]$ 
for each time step t from 0 to T:
    for each state s from 0 to num-states:
        for each s' where edges[s,s'] is a transition probability:
            new-score = viterbi[s,t] * edges[s,s'] * input[s',obs[t]]
            if (viterbi[s',t+1] == 0) or (new-score > viterbi[s', t+1]):
                viterbi[s', t+1] = new-score
                back-pointer[s',t+1] = s
```

3/19/12

19

Viterbi Algorithm

```
# finding the best path
best-final-score = best-final-state = 0
for each state s from 0 to num-states:
    if (viterbi[s,T+1] > best-final-score):
        best-final-state = s
        best-final-score = viterbi[s,T+1]
# start with the last state in the sequence
x = best-final-state
state-sequence.push(x)
for t from T+1 downto 0:
    state-sequence.push(back-pointer[x,t])
    x = back-pointer[x,t]
return state-sequence
```

3/19/12

20