CMPT 825
Natural Language Processing

Anoop Sarkar

http://www.cs.sfu.ca/~anoop

Minimum Cost Edit Distance

- String edit distance: what is the minimum number of changes (char insertions or deletions) to transform the string *intention* into *execution* ?
- Assume cost of insertion is 1 and cost of deletion is 1

Levenshtein Distance

- Cost is fixed across characters
  - Insertion cost is 1
  - Deletion cost is 1
- Two different costs for substitutions
  - Substitution cost is 1 (transformation)
  - Substitution cost is 2 (one deletion + one insertion)

Edit Distance

- Think of it as an alignment between target and source

$$t_1, t_2, \ldots, \ldots, t_n$$
$$s_1, s_2, \ldots, s_m$$

Find *D(n,m)* recursively

$$D(i,j) = min \begin{cases} D(i-1,j) & +\text{cost}(t_i, \emptyset) \textit{ insertion into target} \\ D(i-1,j-1) & +\text{cost}(t_i, s_j) \textit{substitution/identity} \\ D(i,j-1) & +\text{cost}(\emptyset, s_j) \textit{deletion from source} \end{cases}$$

$$D(0,0) = 0$$

$$D(i,0) = D(i-1,0) + \text{cost}(t_i, \emptyset)$$
$$D(0,j) = D(0,j-1) + \text{cost}(\emptyset, s_j)$$

```
Function MinEditDistance (target, source)

 n = length(target)
 m = length(source)
 Create matrix D of size (n+1,m+1)
 D[0,0] = 0

 for i = 1 to n
   D[i,0] = D[i-1,0] + insert-cost

 for j = 1 to m
   D[0,j] = D[0,j-1] + delete-cost

 for i = 1 to n
   for j = 1 to m
     D[i,j] = MIN(D[i-1,j] + insert-cost,
                  D[i-1,j-1] + subst/eq-cost,
                  D[i,j-1] + delete-cost)
 return D[n,m]
```

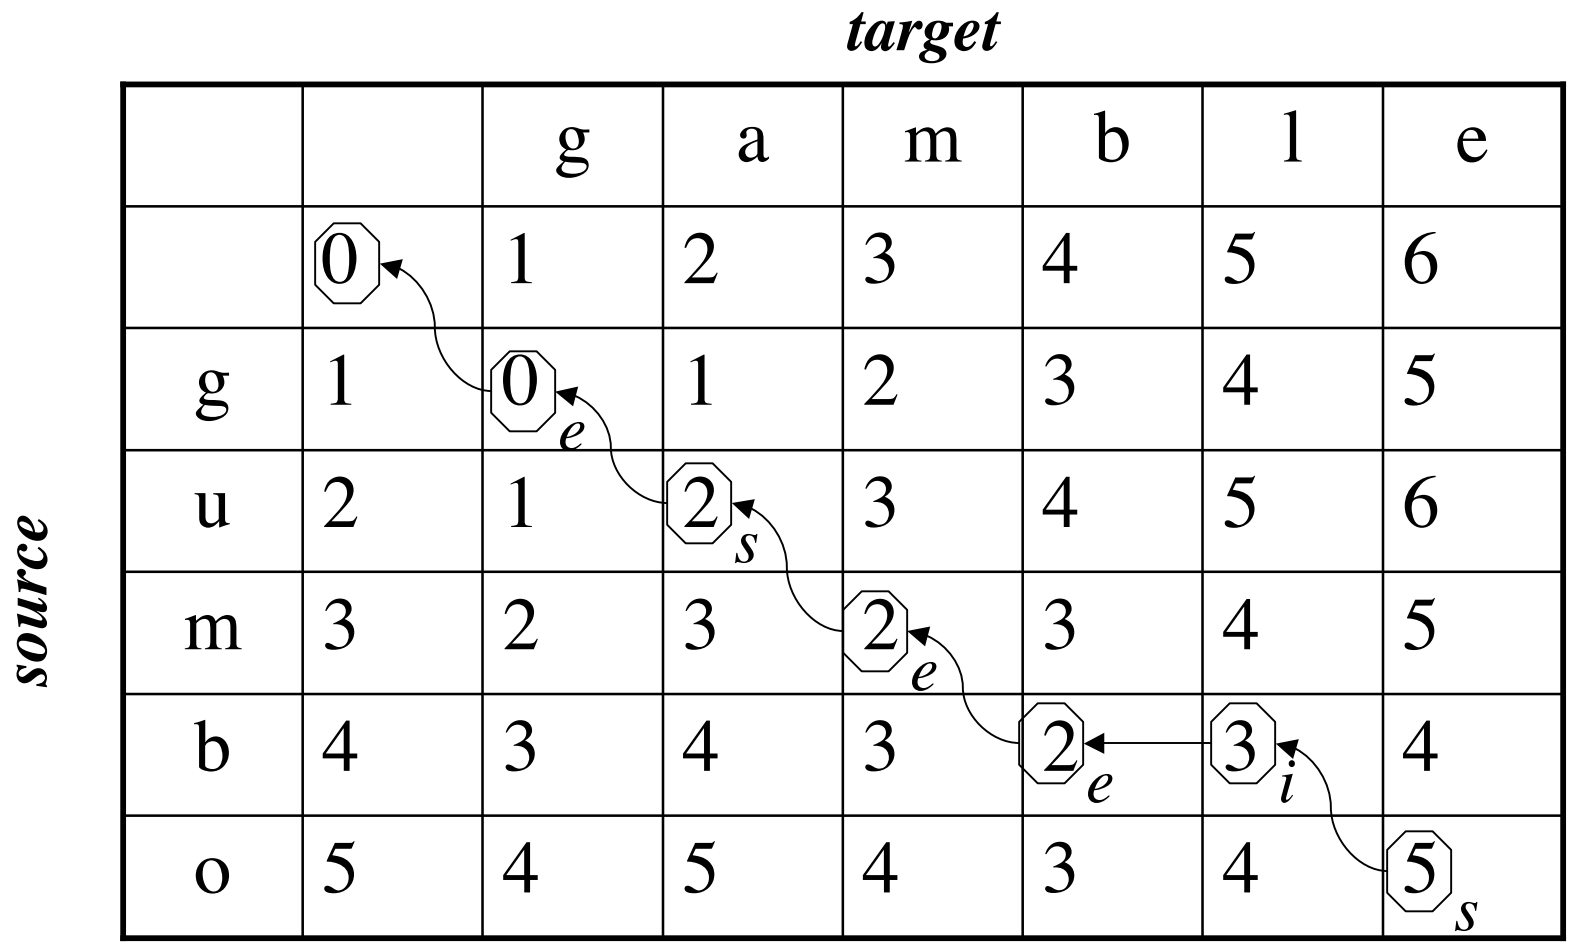
target
source
g a m b l e
0 1 2 3 4 5 6
g 1 0 1 2 3 4 5
u 2 1 2 3 4 5 6
m 3 2 3 2 3 4 5
b 4 3 4 3 2 3 4
o 5 4 5 4 3 4 5
e
s
e
e
i
s

Edit distance

- Useful in many NLP applications
- Comparing system output with human output, e.g. *input:* ibm *output:* IBM vs. Ibm
- Error correction
- Defined over character edits or word edits, e.g. MT evaluation:
  - Foreign investment in Jiangsu ‘s agriculture on the increase
  - Foreign investment in Jiangsu agricultural investment increased

Spelling Correction

- Types of spelling correction
  - non-word error detection

    e.g. *hte* for *the*
  - isolated word error detection

    e.g. *acres* vs. *access* (cannot decide if it is the right word for the context)
  - context-dependent error detection (real world errors)

    e.g. *she is a talented acres* vs. *she is a talented actress*

Noisy Channel Model

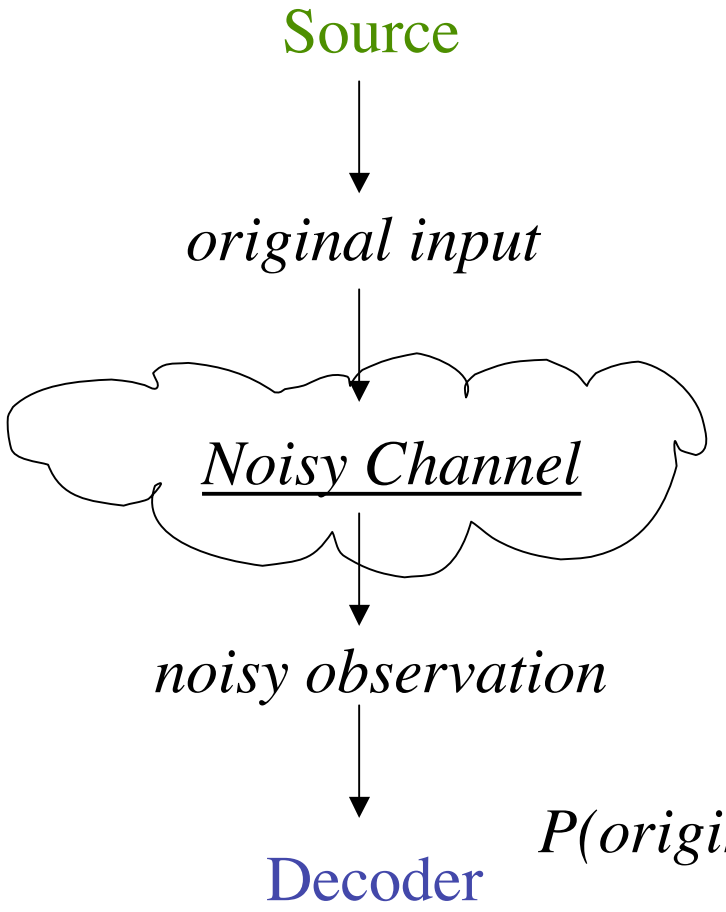

P(original input | noisy obs)

Bayes Rule: *computing P(orig | noisy)*

- let x = *original input*, y = *noisy observation*

$$p(x \mid y) = \frac{p(x,y)}{p(y)} \qquad p(y \mid x) = \frac{p(y,x)}{p(x)}$$

$$p(x,y) = p(y,x)$$

$$p(x \mid y) \times p(y) = p(y \mid x) \times p(x)$$

$$p(x \mid y) = \frac{p(y \mid x) \times p(x)}{\cancel{p(y)}}$$

Bayes Rule

Chain Rule

$$p(a,b,c \mid d) = p(a \mid b,c,d) \times p(b \mid c,d) \times p(c \mid d)$$

Approximations: Bias vs. Variance

$$p(a \mid b,c,d) \approx p(a \mid b,c) \quad \textit{less } \textbf{\textit{bias}}$$
$$p(a \mid b)$$
$$p(a) \quad \textit{more } \textbf{\textit{variance}}$$

Single Error Spelling Correction

- Insertion (addition)
  - acress vs. cress
- Deletion
  - acress vs. actress
- Substitution
  - acress vs. access
- Transposition (reversal)
  - acress vs. caress

Noisy Channel Model for Spelling Correction (Kernighan, Church and Gale, 1990)

- *t* is the typo and *c* is the correct word

$$P(c \mid t) = p(t \mid c) \times p(c)$$

- Find the best candidate for the correct word

$$\hat{c} = \underset{c \in C}{\arg\max} P(t \mid c) \times P(c)$$

$$P(t \mid c) = ?? \qquad P(c) = \frac{f(c)}{N}$$

Noisy Channel Model for Spelling Correction (Kernighan, Church and Gale, 1990) single error, condition on previous letter

$$P(t \mid c) = \begin{cases} \frac{del[c_{p-1}, c_p]}{chars[c_{p-1}, c_p]} & (xy)_c \text{ typed as } (x)_t \\ \frac{ins[c_{p-1}, t_p]}{chars[c_{p-1}]} & (x)_c \text{ typed as } (xy)_t \\ \frac{sub[t_p, c_p]}{chars[c_p]} & (y)_c \text{ typed as } (x)_t \\ \frac{rev[c_p, c_{p+1}]}{chars[c_p, c_{p+1}]} & (xy)_c \text{ typed as } (yx)_t \end{cases}$$

Noisy Channel model for Spelling Correction

- The *del*, *ins*, *sub*, *rev* matrix values need data in which contain known errors (**training data**)
- Accuracy on single errors on unseen data (**test data**)

Noisy Channel model for Spelling Correction

- Experiments: 87% accuracy for machine vs. 98% average human accuracy
- What are the limitations of this model?

*… was called a "stellar and versatile **acress** whose combination of sass and glamour has defined her …*

KCG model best guess is **acres**